AI FOR DATA ANALYSIS · LESSON 1 OF 4

Generative AI for Analysts: How It Works, Prompting & AI vs GenAI vs Agentic AI

Understand what generative AI is and how large language models behave (tokens, context, temperature, hallucination), how AI, machine learning, deep learning, GenAI and agentic AI differ, and how to write prompts that give analysts accurate, checkable results, with reusable templates and privacy rules.

DVA with AI · class 1 | ✦ Models and tools as of 2026 | Beginner | ⏱ 110 min

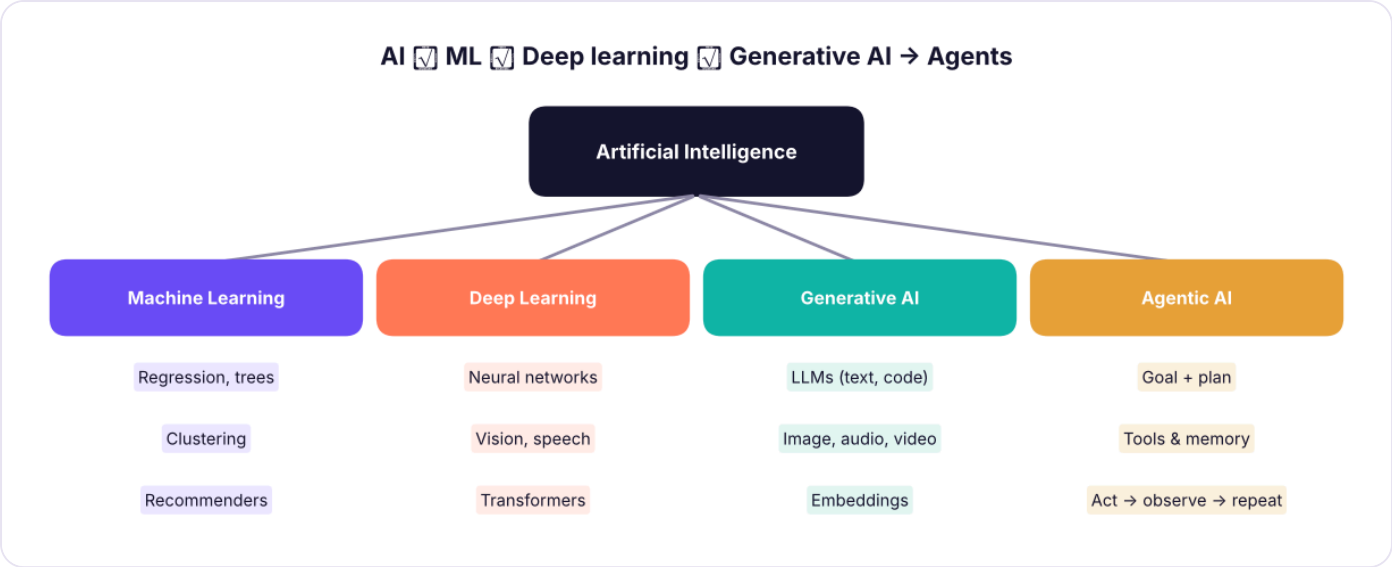
By the end of this lesson you can

- Explain the relationship between AI, ML, deep learning, generative AI and agentic AI
- Describe in plain language how an LLM generates text and why it can be confidently wrong (hallucination)
- Write prompts using role, context, task, format and constraints, with examples and step-by-step reasoning
- Use prompt templates for common analyst tasks and know what never to paste into an AI tool
- Choose a suitable tool and check an AI answer before using it

Datasets used

`superstore_sales.csv`

The AI family tree



How the terms relate. Each ring is a subset of the one before it. Agentic AI uses generative models as its brain but adds goals, tools and loops.

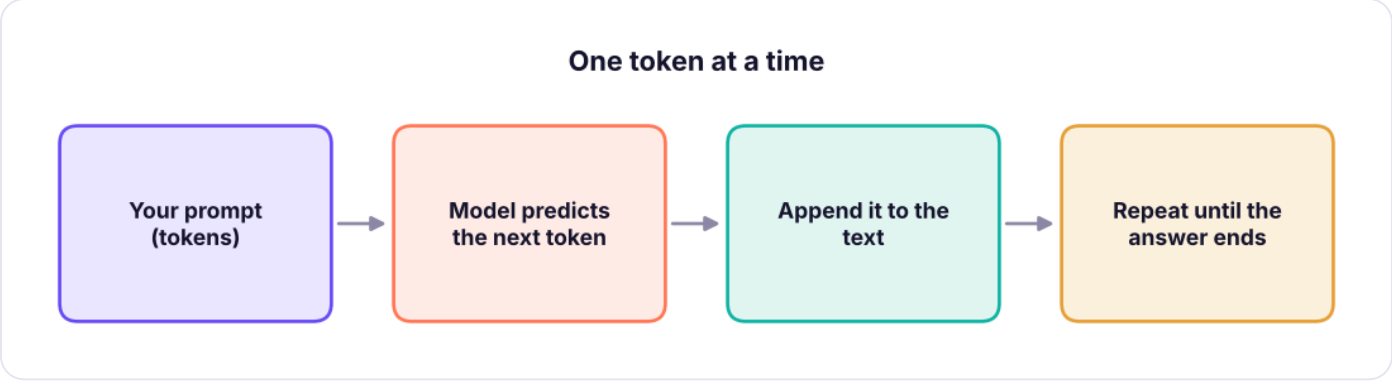
Term	What it means	Example
AI	Machines doing tasks that normally need human intelligence	Rule-based chess engine, spam filter
Machine learning (ML)	Systems that **learn patterns from data** instead of following hand-written rules	Predict loan default from past loans
Deep learning	ML with many-layered neural networks	Image recognition, speech-to-text
Generative AI (GenAI)	Models that **create new content** (text, images, code, audio)	ChatGPT, Claude, Gemini, Copilot, Midjourney
LLM	Large language model: the GenAI engine for text and code	GPT, Claude, Gemini, Llama
Agentic AI	An LLM **plus goals, tools and a loop**: it plans, calls tools (run SQL, search, edit files), checks results and continues until the goal is met	An analyst agent that explores a database and writes the report
AGI	Hypothetical AI at human level across tasks	Not here yet; ignore the hype

AI vs GenAI vs agentic AI

	Traditional ML	Generative AI	Agentic AI
Output	A prediction (number, label)	New content (text, code, image)	A completed multi-step task
Input	Structured data	A prompt (text, files, images)	A goal in plain language
Control	You train and deploy a model	You ask once, get one answer	The system decides the next step
Example	"Will this customer churn?" → 0.82	"Summarise these reviews" → paragraph	"Find why sales fell in Q3 and draft a memo" → runs queries, charts, writes
Main risk	Bias, drift	Hallucination	Hallucination **plus actions**: wrong tool calls, runaway cost, security

How an LLM works (without maths)

An LLM is trained on vast amounts of text to predict **the next token** (a word or word piece). Repeating that prediction one token at a time produces sentences, code and tables. It has learnt statistical patterns of language, facts and reasoning styles, but it does **not look things up** unless it is given a tool or documents.



Generating an answer: the prompt is split into tokens; the model repeatedly predicts the next token from everything so far.

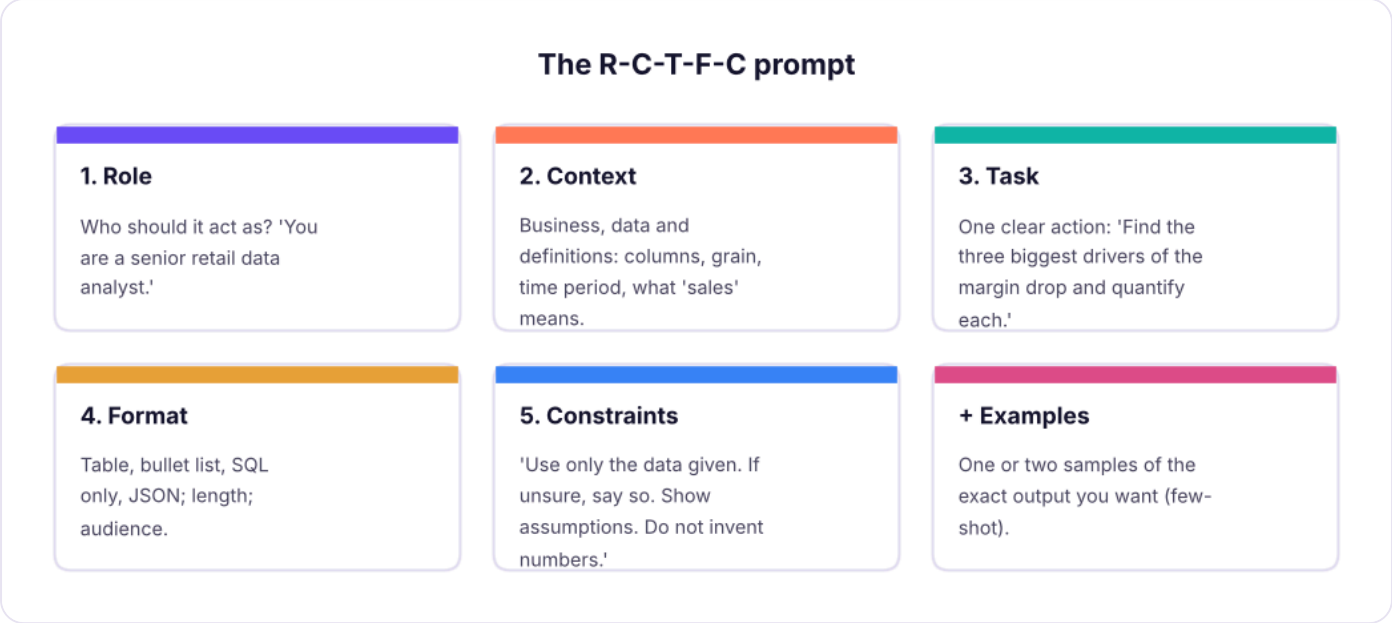
Concept	Meaning	Why analysts care
Token	Roughly ¾ of an English word (about 4 characters)	Limits and costs are measured in tokens
Context window	How much text the model can consider at once (prompt + answer)	A long file may not fit; the model forgets what is outside
Temperature	Randomness of word choice	Low (0–0.3) for SQL, maths and extraction; higher for brainstorming
System prompt	Hidden instructions that set behaviour	Custom GPTs/Projects store your standing instructions
Training cut-off	The model's knowledge stops at some date	It may not know recent tools or prices; give it documents
Grounding / RAG	Supplying real documents or data in the prompt (Retrieval-Augmented Generation)	Cuts hallucination: "answer only from this file"
Tool use / function calling	The model asks the app to run a calculator, database query or code	Makes answers computed instead of guessed
Multimodal	Accepts images, PDFs, audio as well as text	Paste a chart screenshot and ask what it shows

WATCH OUT **Hallucination**

Because an LLM generates *plausible* text, it can state false things with complete confidence: invented statistics, non-existent functions, wrong column names, fake citations. **Never use an AI answer as a fact until you have checked it**: run the code, recompute the number, open the source.

Prompt engineering for analysts

A prompt is an instruction plus the material the model needs. The most reliable structure has five parts:



A prompt recipe: Role, Context, Task, Format, Constraints.

Weak vs strong prompt

Weak: *"Analyse my sales data."*

Strong:

Text

```
Role: You are a senior retail data analyst for an Indian omnichannel retailer.
Context: Table superstore(order_id, order_date, region, category, sub_category, sales,
discount, profit),
         one row per order line, Jan 2023–Dec 2025, amounts in rupees. Margin = profit /
sales.
Task: Explain why overall margin is only about 9%. Look at category, sub-category and
discount band.
Format: (1) a 3-bullet summary for the CFO, (2) a table of the 5 worst sub-categories
(sales, profit, margin),
        (3) the SQL Server queries you used.
Constraints: Use only the columns listed. State any assumption. If a number cannot be
computed from the table, say so.
```

The strong prompt names the grain, the metric definition, the output format and the rules for uncertainty. The result is far more likely to be right *and* checkable.

Core techniques

Technique	How	Use for
Be specific	Name columns, periods, units, audience	Everything
Few-shot examples	Show 2–3 input→output pairs	Classification, extraction, tone matching
Step by step	"First list the assumptions, then compute, then summarise"	Multi-step analysis; makes errors visible
Ask for checks	"Then list three ways this answer could be wrong"	Self-critique
Structured output	"Return JSON with keys..."	Feeding results into code
Delimiters	Put data in `"""triple quotes"""` or XML tags	Stops the model confusing instructions with data
Iterate	Treat the first answer as a draft; refine with follow-ups	Always
Give it the schema, not the data	Column names and 5 sample rows, not the whole file	Privacy, token limits

Prompt templates for common analyst jobs

Job	Template (replace brackets)
Explain a formula	"Explain this Excel formula step by step for a beginner and give a simpler alternative: [formula]"
Write SQL	"Dialect: SQL Server. Tables: [schema with keys and grain]. Write a query that [question]. Return only SQL, then list assumptions."
Debug code/SQL	"This [language] query gives [wrong result / this error]. Expected: [x]. Query: [code]. Find the cause and show the fix."
Clean data	"Here are 10 sample values of column [x]: [list]. Propose normalisation rules and a mapping table; flag ambiguous cases."
Interpret results	"Here is a result table: [table]. Write 3 insights for a non-technical manager, each with a number, and 2 caveats."
Choose a chart	"I want to show [message] to [audience] with [data]. Recommend a chart and what to label."
Draft a report	"Turn these findings into a one-page memo with: situation, key findings, recommendation, risks. Keep it under 250 words."
Interview practice	"Act as an interviewer for a data analyst role; ask one SQL question at a time and give feedback on my answer."

EXAMPLE

From vague to checkable

Prompt: *"Explain this formula:* `=IFERROR(INDEX($C:$C,MATCH(1,($A:$A=F2)*($B:$B=G2),0)),"Not found")` *"* A good answer says it looks up C where A equals F2 **and** B equals G2 (a two-condition lookup), returns "Not found" if none, and notes it needs array evaluation in older Excel. You can verify by testing it on 3 rows.

Choosing and using AI tools

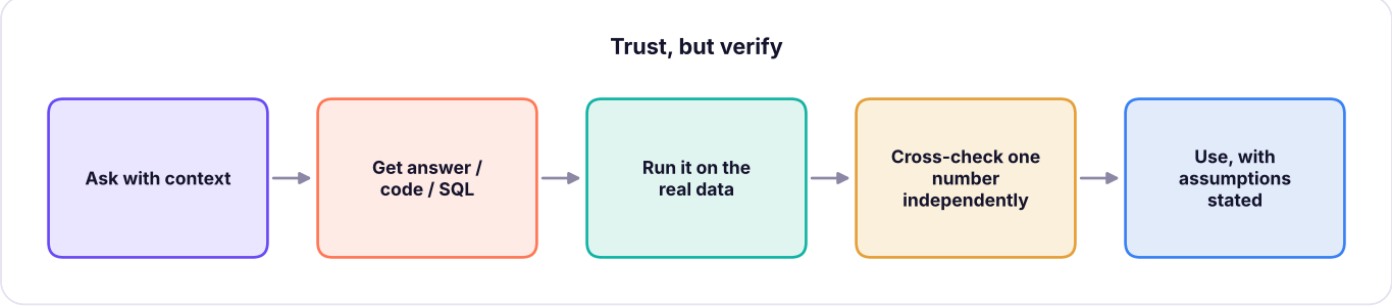
Tool type	Examples	Good for	Watch out
General chat assistants	Claude, ChatGPT, Gemini, Copilot, Grok, Perplexity	Explaining, drafting, code, brainstorming	Data privacy; hallucination
Search-grounded	Perplexity, Gemini with search, Copilot with web	Fresh facts with citations	Check the cited pages; sources vary in quality
Document/notebook tools	NotebookLM, Projects, Custom GPTs	Answering **only from your files**	Quality depends on the files
Embedded in apps	Copilot in Excel/Power BI, Gemini in Sheets, Metabase/ MinusX, Claude in Excel	Work where the data is	Permissions; see next lessons
Coding assistants and agents	Claude Code, Cursor, GitHub Copilot, Codex	Writing and running code, analysis scripts	Review diffs; run tests; secrets

Never paste

1. **Personal data** (names with phone numbers, emails, Aadhaar/PAN, health, salary of identifiable people) unless your company's approved tool and policy allow it.
2. **Confidential business data** (unreleased financials, customer lists, contracts, source code, credentials, API keys) into public tools.
3. **Anything you are not allowed to share**: client data under NDA.

Use **enterprise versions** with data-protection terms, **mask or sample** the data (replace names with IDs, share the schema and 5 fake rows), and follow your organisation's AI policy. In India, personal data is also covered by the Digital Personal Data Protection Act, 2023.

A verification habit for every AI answer



The analyst's rule: generate with AI, verify with data.

Here is why verification matters. Suppose an AI writes the following query for "average order value per region". It looks right, runs without error and returns plausible numbers, but it averages order **lines**, not **orders**:

Python

```
import duckdb, pandas as pd
con = duckdb.connect()
con.execute("CREATE TABLE superstore AS SELECT * FROM
read_csv_auto('superstore_sales.csv')")
wrong = con.execute("SELECT region, ROUND(AVG(sales), 0) AS avg_order_value FROM superstore
GROUP BY region ORDER BY region").df()
right = con.execute("""SELECT region, ROUND(AVG(order_total), 0) AS avg_order_value
                       FROM (SELECT region, order_id, SUM(sales) AS order_total FROM
superstore GROUP BY region, order_id) GROUP BY region ORDER BY region""").df()
pd.DataFrame({"region": wrong.region, "AI query (per line)": wrong.avg_order_value,
"Correct (per order)": right.avg_order_value})
```

OUTPUT

	region	AI query (per line)	Correct (per order)
0	East	30,521	60,684
1	North	32,038	61,892
2	South	32,580	64,925
3	West	32,192	64,067

4 rows × 3 columns

The AI's version understates order value by roughly half, because an average order has about two lines. A one-line sanity check (total sales ÷ number of orders) would have caught it:

Python

```
tot = con.execute("SELECT SUM(sales) / COUNT(DISTINCT order_id) FROM
superstore").fetchone()[0]
print(f"Total sales / distinct orders = {tot:,.0f}")
```

OUTPUT

```
Total sales / distinct orders = 62,942
```

COMMON MISTAKES **Common AI-assistant failure modes**

1. **Wrong grain**: counting lines as orders, double-counting after a join (Lesson 8 in SQL).
2. **Invented columns or functions**: the code references a field that does not exist.
3. **Silent assumptions**: it picks a date range, currency or definition without saying.
4. **Plausible but unverified numbers**: especially when asked to "estimate" or "recall" data it cannot see.
5. **Over-agreeing**: if you assert something wrong, it may follow you. Ask it to challenge your assumption.
6. **Stale knowledge**: new tool versions, prices, laws.

Responsible use

- **Accountability stays with you**: you sign off on the analysis, not the model.
- **Bias**: models mirror training data; check outputs for unfair assumptions about people or regions.
- **Transparency**: tell stakeholders when AI helped, and keep the prompt and verification notes.
- **Copyright and attribution**: do not present generated text or images as someone else's work; check licences.
- **Security**: treat prompts and uploaded files as data leaving your machine; beware prompt injection when an AI reads untrusted documents or web pages.

DO IT WITH AI **Do it with AI**

Pick a task you did earlier in this course (for example the SQL discount analysis). Ask an AI assistant to produce the answer from the schema, run its SQL, and compare with the result in this handbook. Where it differs, work out *why* (grain? date range? definition?). This single exercise teaches more about AI limits than any lecture.

Summary

KEY TAKEAWAYS

- AI ⊃ ML ⊃ deep learning ⊃ generative AI; **agentic AI** = an LLM with goals, tools and a loop.
- LLMs predict the next token; they do not look things up unless given documents or tools. They can **hallucinate**, so verify every answer.
- Good prompts have **role, context, task, format and constraints**, use examples and step-by-step reasoning, and give the schema (not the sensitive data).
- Use templates for SQL, formulas, cleaning, insights, memos and interview practice; iterate.
- Never paste personal or confidential data into public tools; mask or sample, and follow your company policy.
- The analyst's loop: ask with context → run on real data → cross-check one number → state assumptions.

Quick quiz

6 questions · tap an answer to check it

1. Which statement best describes agentic AI?

A A model that only classifies images

B An LLM-based system that plans, uses tools and repeats steps to achieve a goal

C Any chatbot

D A rule-based macro

Answer: B. Agents add goals, tools and a decision loop on top of a language model.

2. Why can an LLM state wrong facts confidently?

A It lacks memory

B It generates plausible text from patterns and does not verify facts unless given sources or tools

C It is connected to the internet

D It always copies its training data

Answer: B. Fluency is not accuracy.

3. Which setting is best for generating SQL?

A High temperature

B Low temperature

C No context

D Maximum length

Answer: B. Low temperature reduces random variation for precise tasks.

4. What makes a prompt more reliable for analysis?

A Shorter wording

B Role, context (schema and definitions), a clear task, an output format and constraints

C Asking for "everything"

D Hiding the table names

Answer: B. Specifics reduce guesswork and make the answer checkable.

5. A teammate wants to paste a customer list with phone numbers into a public chatbot. Best advice?

- **A** Go ahead
- **B** Do not; mask or sample the data, or use an approved enterprise tool
- **C** Delete the email column only
- **D** Convert it to PDF first

Answer: B. Personal data should not go into tools not approved for it.

6. An AI query returns an "average order value" half the size of total sales ÷ orders. The likely error is:

- **A** Wrong chart
- **B** Averaging order lines instead of order totals (wrong grain)
- **C** Too many decimals
- **D** A missing index

Answer: B. Always sanity-check results against a simple independent calculation.

EXERCISE **Prompt rewrite**

Rewrite this weak prompt into a strong one and list three checks you will run on the answer: *"Tell me about churn in my telecom data."* The data is `telecom_churn.csv` with columns customer_id, senior_citizen, tenure_months, contract, internet_service, payment_method, monthly_charges, total_charges, support_calls, churn (Yes/No).

Show solution (worked answer)

Strong prompt: "Role: senior telecom data analyst. Context: table telecom_churn(...columns...), one row per subscriber, churn = Yes/No. Task: find which contract types, tenure bands and support-call counts have the highest churn rate and quantify each. Format: a table (group, customers, churners, churn %), three bullets for a retention manager, and the SQL/pandas code used. Constraints: use only these columns, state assumptions, do not infer causes you cannot test, flag small groups (< 50 customers)." Checks: (1) recompute the overall churn rate by hand (count of Yes ÷ rows) and match it; (2) confirm each group's customer counts add up to the total; (3) spot-check one group with an independent query.

INTERVIEW CORNER **Questions you may be asked**

Q: Difference between AI, ML and generative AI? AI is the broad field; ML learns patterns from data to predict or classify; generative AI (a deep-learning subset) creates new content such as text, code and images.

Q: What is a hallucination and how do you reduce it? Confident but false output. Reduce it by grounding answers in provided documents or data (RAG), using tools that compute results, lowering temperature, asking for sources and assumptions, and verifying independently.

Q: How do you write a good prompt for a data task? State the role, context (schema, grain, definitions), a single clear task, the required output format, and constraints (use only given data, state assumptions); give examples; iterate.

Q: What is agentic AI? Systems where an LLM plans and executes multi-step tasks using tools, observes results and continues until a goal is met, with human oversight for risky actions.

Practise more online

Free, trusted places to test yourself on this topic. Links open the original site in a new tab.